Everyone Knows Trump's Tweets Move Markets.

I Measured It: the Connection Is Real — the Direction Isn't.

Nebius Serverless AI Builders Challenge · #NebiusServerlessChallenge · Koral Zakai

The question everyone already answers

Every trading desk has the anecdote: *"he tweeted about Intel and the stock ran for months."* I built a research pipeline on **Nebius Serverless AI** to test it properly: 78,130 public posts, 62 tickers of committed daily bars, and **Llama-3.3-70B** on Nebius AI Studio reading the text of every market-relevant tweet — no fine-tuning, no market data in the prompt. 476 tweets, 1,296 instrument calls, every test pre-registered before scoring.

What the data actually says

The connection is real — but it runs backwards. The market moves first; then he posts.
Forward, the tradeable direction, is a coin flip: 0 of 63 pre-registered tests survive.

Intel had already gained **+105% in 21 sessions** when he posted the famous "Intel Stock continues to rise." He tweeted about Tesla **thirteen times** — while it fell 26%. His posts land mid-trend, so memory reorders tweet-and-move into tweet-then-move. The observed correlation is genuine; the implied causation is backwards. Three mechanisms make the anecdotes feel true: **reverse causality** (he narrates moves that already happened), **mean reversion** (big oil moves recover identically with no tweet at all), and **saturation** (he posts on 83% of trading sessions, so every big move has a same-day tweet available to blame).

what we tested	result	verdict
Daily directional calls (held-out test)	52.2% hit rate (n=67, p=0.40)	coin flip
1,296 instrument calls vs the S&P; 500	48.1 / 47.3 / 48.6 / 49.2 % at 1d/3d/1w/1mo	coin flip
Pre-registered event study, 63 cells	0 survive BH correction (min p = 0.76)	null
Intraday shock study, 30 cells	0 survive (detects effects 5x smaller)	null
"When can we trust it?" meta-model	AUC 0.593 validation -> 0.431 unseen	worse than chance
The mirror test (market -> his feed)	posting intensity tracks the move behind it	suggestive only

The mirror test is held to the same bar that killed the forward signal: its raw p=0.033 fails our own multiple-comparison correction, so it is reported as suggestive — never as a survivor.

Seven beautiful, publishable, completely false findings

The real engineering content. Seven times this pipeline produced exactly the discovery the anecdotes promise — statistically significant, story-shaped, and wrong. Each was caught by a control, not by luck:

1	**The tie-break.** Baskets splitting 1-of-2 were counted as full WINS. That single ">= 0.5" produced the headline: 61.8%, p=0.017.	**0.618 -> 0.507**
2	**The null that could not say no.** He posts on 83% of sessions; sampling a null pool ~ the whole population collapsed its variance and floored every p-value.	**22 hits -> 0**
3	**The same-bar burst.** Five tweets in ten minutes, one hourly bar: one event counted five times. Caught because a "renovated Palm Room" tweet topped the movers list.	**9 cells -> 0**
4	**The frozen clock.** US open hardcoded to 13:30 UTC; in winter NYSE opens 14:30. Every winter morning tweet was scored against the wrong session.	**8 rows re-anchored**
5	**DJT is his signature.** He signs posts "President DJT". The ticker rule read that as a mention of his own listed company.	**754/762 fake**
6	**Negated-return math.** -E(-X) as a backward return is fine for small moves; on a stock that had doubled it reported +44.8% for a +105.3% run-up.	**+44.8 -> +105.3%**
7	**"intel" is not Intel.** In a political corpus "intel" means intelligence: 51 of 55 matches were spy-agency tweets, sitting inside a CORPORATE cohort. The mirror bug hid the 4 counterexamples to our own headline exhibit.	**cohort 43 -> 13**

After every fix the verdict was re-scored. It never changed: nothing survives.

What shipped on Nebius Serverless

Batch Jobs (CPU-only, deterministic, hash-stamped) run the research pipeline and write a validation manifest. A **Serverless Endpoint** serves POST /predict — and it **refuses to boot** if its live prompt hash or corpus hash disagree with that manifest: the Job produces the numbers, the Endpoint may only cite them. Because nothing survived correction, the deployed endpoint **abstains on every tweet**, explains why, and quotes no accuracy — a calibrated refusal is the correct output. A feedback Job scores every served call after its horizon closes, as a prospective replication set with a correction that grows stricter on every look — the loop cannot be re-run into significance.

The serverless layer contains zero science: jobs and serving only marshal I/O around one pure, tested engine. One CPU image builds both, so batch and serve cannot drift.

Why a null result is the win

A pipeline that can only say yes is an anecdote generator with extra steps. This one was built to be able to say **no** — point-in-time feature discipline, purged evaluation, permutation nulls, Benjamini-Hochberg correction over a pre-registered registry — and then it was made to prove it could still say yes (a planted-signal canary must be detected, or the null is meaningless). Six of the seven artifacts were caught by a control. The seventh was caught because a tweet about a renovated Palm Room topped a leaderboard.

Live dashboard — the "Myth-Busting Quantitative Terminal" (interactive, 6 tabs):
koralzakai.github.io/stocksPredictionAfterTweet/dashboard.html
Live endpoint (Nebius Serverless, try any tweet):
port8080-xaqg3twr3yhr61q.tunnel.applications.eu-north1.nebius.cloud/docs
Code (MIT, reproducible for $0 offline — 177 tests):
github.com/KoralZakai/stocksPredictionAfterTweet

CPU-only · public data only · fully deterministic (seeded, hash-stamped) · offline reproduction ~10 s, $0 · research output, not investment advice